

PANDAS

Avijit Bose

Assistant Professor

Department of Computer science & Engineering,
MCKV institute of engineering,Liluah,howrah-711204

Pandas

- Pandas is a python library used for working with data sets.
- It has functions for analyzing, cleaning, exploring and manipulating data.
- Pandas allow us to analyze big data and make conclusions based on statistical theories.
- Pandas can clean messy data sets and make them readable and relevant.
- Relevant data is very important in data science.

Pandas

- Is there a correlation between two or more columns?
- What is average value?
- Max value?
- Min value?

How to install pandas

- `C:\Users\Your Name>pip install pandas`

How to import pandas

- `import pandas`
- `mydataset = {`
- `'cars': ["BMW", "Volvo", "Ford"],`
- `'passings': [3, 7, 2]`
- `}`
- `myvar = pandas.DataFrame(mydataset)`
- `print(myvar)`

Checking for pandas version

- `import pandas as pd`
`print(pd.__version__)`

Pandas Series

Pandas series is like a column of a table

It is a one dimensional array holding data of any type.

```
import pandas as pd
```

```
a = [1, 7, 2]
```

```
myvar = pd.Series(a)
```

```
print(myvar)
```

Labels

If nothing else is specified, the values are labeled with their index number. First value has index 0, second value has index 1 etc.

This label can be used to access a specified value.

```
print(myvar[0])
```


Creating Labels

- With the index argument, you can name your own labels.

```
import pandas as pd
```

```
a = [1, 7, 2]
```

```
myvar = pd.Series(a, index = ["x", "y", "z"])  
print(myvar)
```

When you have created labels, you can access an item by referring to the label.

```
print(myvar["y"])=
```

Key value Objects in Series

- We can also use a key/value object, like a dictionary, when creating a Series.

```
import pandas as pd
```

```
calories = {"day1": 420, "day2": 380, "day3": 390}
```

```
myvar = pd.Series(calories)
```

```
print(myvar)
```

Key value objects in dictionary

- To select only some of the items in the dictionary, use the index argument and specify only the items you want to include in the Series.
- Create a Series using only data from "day1" and "day2":

```
import pandas as pd
```

```
    calories = {"day1": 420, "day2": 380, "day3": 390}
```

```
    myvar = pd.Series(calories, index = ["day1",  
"day2"])
```

```
    print(myvar)
```

Data Frames

- Data sets in Pandas are usually multi-dimensional tables, called DataFrames.
- Series is like a column, a DataFrame is the whole table.
- Create a DataFrame from two Series:

```
import pandas as pd
```

```
    data = {  
        "calories": [420, 380, 390],  
        "duration": [50, 40, 45]  
    }
```

```
myvar = pd.DataFrame(data)
```

```
print(myvar)
```

What is a data frame

- A Pandas DataFrame is a 2 dimensional data structure, like a 2 dimensional array, or a table with rows and columns.
- Create a simple Pandas DataFrame:

```
import pandas as pd
```

```
data = {  
    "calories": [420, 380, 390],  
    "duration": [50, 40, 45]  
}
```

```
#load data into a DataFrame object:  
df = pd.DataFrame(data)
```

```
print(df)
```

Locate Row

- As we can see from the result above, the DataFrame is like a table with rows and columns.
- Pandas use the loc attribute to return one or more specified row(s)

Return row 0:

#refer to the row index:

```
print(df.loc[0])
```

Locate row

- Example
- Return row 0 and 1:

#use a list of indexes:

```
print(df.loc[[0, 1]])
```

Locating named indexes

- Named Indexes
- With the index argument, you can name your own indexes.
- Add a list of names to give each row a name:

```
import pandas as pd
```

```
    data = {  
        "calories": [420, 380, 390],  
        "duration": [50, 40, 45]  
    }
```

```
    df = pd.DataFrame(data, index = ["day1", "day2", "day3"])  
    print(df)
```


Locating Named indexes

- Use the named index in the loc attribute to return the specified row(s).
- Return "day2":

#refer to the named index:

```
print(df.loc["day2"])
```

Loading files into a data frame

- Load a comma separated file (CSV file) into a DataFrame:
- ```
import pandas as pd
df = pd.read_csv('data.csv')
print(df)
```

# Read CSV Files

- A simple way to store big data sets is to use CSV files (comma separated files).
- CSV files contains plain text and is a well know format that can be read by everyone including Pandas.
- In our examples we will be using a CSV file called 'data.csv'.

# Read csv file

- Load the CSV into a DataFrame:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
print(df.to_string())
```

# Read csv file

- By default, when you print a DataFrame, you will only get the first 5 rows, and the last 5 rows:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
print(df) # if data set is huge then only first 5 and
last 5
```

# Reading JSON File

- Big data sets are often stored, or extracted as JSON.
- JSON is plain text, but has the format of an object, and is well known in the world of programming, including Pandas.
- In our examples we will be using a JSON file called 'avi.json'.

# Reading JSON File

```
import pandas as pd
```

```
df = pd.read_json('data.json')
```

```
print(df.to_string())
```

# Dictionary as JSON

- JSON = Python Dictionary
- JSON objects have the same format as Python dictionaries.



# Example

```
import pandas as pd
```

```
data = {
 "Duration":{
 "0":60,
 "1":60,
 "2":60,
 "3":45,
 "4":45,
 "5":60
 },
 "Pulse":{
 "0":110,
 "1":117,
 "2":103,
 "3":109,
 "4":117,
 "5":102
 },
}
```

# Example

```
"Maxpulse":{
 "0":130,
 "1":145,
 "2":135,
 "3":175,
 "4":148,
 "5":127
},
"Calories":{
 "0":409,
 "1":479,
 "2":340,
 "3":282,
 "4":406,
 "5":300
}
}
```

# Example

```
df = pd.DataFrame(data)
print(df)
```

```
/* Printing the first 10 rows of data set*/
import pandas as pd
df = pd.read_csv('data.csv')
print(df.head(10))
```

```
/* Print the first 5 rows of data */
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
print(df.head())
```

# Printing the data frame

Printing the last 5 rows of data frame

```
print(df.tail())
```

```
/* Info about the data set */
```

In order to print the information of the data we have to give the command

```
print(df.info())
```

The result tells us there are n rows and m columns:

And the name of each column, with the data type:

# Concept of Null Values

- The info() method also tells us how many Non-Null values there are present in each column, and in our data set it seems like there are 164 of 169 Non-Null values in the "Calories" column.
- Which means that there are 5 rows with no value at all, in the "Calories" column, for whatever reason.
- Empty values, or Null values, can be bad when analyzing data, and you should consider removing rows with empty values. This is a step towards what is called cleaning data,

# Bad data – Data Cleaning

Bad data could be

- (a) empty cells
- (b) data in wrong format
- (c) wrong data
- (d) duplicates

# Return a new data frame with no empty cells

- Return a new data frame with no empty cells

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
new_df = df.dropna()
```

```
print(new_df.to_string())
```

# What is the purpose of dropna

- `dropna()` method returns a new data frame and will not change the original.
- If we want to change the original DataFrame we use the `inplace= True` assignment,.



# Example

- Remove all rows with NULL values:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
df.dropna(inplace = True)
```

```
print(df.to_string())
```

Now, the `dropna(inplace = True)` will NOT return a new DataFrame, but it will remove all rows containing NULL values from the original DataFrame.

# Replace empty values

- Another way of dealing with empty cells is to insert a *new* value instead.
- This way you do not have to delete entire rows just because of some empty cells.
- The `fillna()` method allows us to replace empty cells with a value:

# Example

- Replace NULL values with the number 130:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
df.fillna(130, inplace = True)
```

# Replace Only for a specified column

- To only replace empty values for one column, specify the *column name* for the Data Frame:
- Replace NULL values in the "Calories" columns with the number 130:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
df["Calories"].fillna(130, inplace = True)
```

# Replace Using Mean,Median,Mode

- A common way to replace empty cells, is to calculate the mean, median or mode value of the column.
- Pandas uses the `mean()` `median()` and `mode()` methods to calculate the respective values for a specified column:

# Replace Using Mean, Median, or Mode

- Example
- Calculate the MEAN, and replace any empty values with it:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
x = df["Calories"].mean()
```

```
df["Calories"].fillna(x, inplace = True)
```

# Replace Using Mean, Median, or Mode

- Calculate the MEDIAN, and replace any empty values with it:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
x = df["Calories"].median()
```

```
df["Calories"].fillna(x, inplace = True)
```

# Replace Using Mean, Median, or Mode

- Calculate the MODE, and replace any empty values with it:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
x = df["Calories"].mode()[0]
```

```
df["Calories"].fillna(x, inplace = True)
```



# Convert Into a Correct Format

- In our Data Frame, we have two cells with the wrong format. Check out row 22 and 26, the 'Date' column should be a string that represents a date:

# Convert Into a Correct Format

- Let's try to convert all cells in the 'Date' column into dates.
- Pandas has a `to_datetime()` method for this:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
df['Date'] = pd.to_datetime(df['Date'])
```

```
print(df.to_string())
```

# What will be the problem

- As you can see from the result, the date in row 26 where fixed, but the empty date in row 22 got a NaT (Not a Time) value, in other words an empty value. One way to deal with empty values is simply removing the entire row.
- Remove rows with a NULL value in the "Date" column:
- `df.dropna(subset=['Date'], inplace = True)`

# Pandas- Fixing wrong data

- Wrong data can be anything for example the data in place of 1.78 may be 178.
- Some of the data may be out of range then also the problem needs to be resolved.
- For example we can write something like :-
- `df.loc[7,'duration']=45`
- For small data sets we can have values checked but what about large data sets.

# Large data sets

```
for x in df.index:
```

```
 if df.loc[x, "Duration"] > 120:
```

```
 df.loc[x, "Duration"] = 120
```

# Removing Rows

- Another way of handling wrong data is to remove the rows that contain the wrong data
- Delete rows where duration is higher than 120

for x in df.index:

if df.loc[x, "Duration"] > 120:

df.drop(x, inplace = True)

Inplace will make permanent change in df.

# Discovering and removing duplicates

- Duplicate rows are rows that have been registered more than one time.
- To discover duplicates we can use the duplicated method.
- The duplicated () method returns a boolean value for each row.
- Returns true for every row that is a duplicate , otherwise false.
- `print(df.duplicated())`

# Removing duplicates

- To remove duplicates we use the `drop_duplicates()` method
- `df.drop_duplicates(inplace = True)`
- The `(inplace = True)` will make sure that the method does NOT return a new DataFrame, but it will remove all duplicates from the original DataFrame.



# Pandas data

- A great aspect of the Pandas module is the `corr()` method.
- The `corr()` method calculates the relationship between each column in your data set.
- `df.corr()`

# Pandas data

- The Result of the `corr()` method is a table with a lot of numbers that represents how well the relationship is between two columns.
- The number varies from -1 to 1.
- 1 means that there is a 1 to 1 relationship (a perfect correlation), and for this data set, each time a value went up in the first column, the other one went up as well.
- 0.9 is also a good relationship, and if you increase one value, the other will probably increase as well.
- -0.9 would be just as good relationship as 0.9, but if you increase one value, the other will probably go down.
- 0.2 means NOT a good relationship, meaning that if one value goes up does not mean that the other will.

# Pandas data

- Perfect Correlation:

We can see that "Duration" and "Duration" got the number 1.000000, which makes sense, each column always has a perfect relationship with itself.

- Good Correlation:

"Duration" and "Calories" got a 0.922721 correlation, which is a very good correlation, and we can predict that the longer you work out, the more calories you burn, and the other way around: if you burned a lot of calories, you probably had a long work out.

- Bad Correlation:

"Duration" and "Maxpulse" got a 0.009403 correlation, which is a very bad correlation, meaning that we can not predict the max pulse by just looking at the duration of the work out, and vice versa.